



CXL Type-2 calibration — device model vs. characterization rails vs. real host

Device: patched cxl-type2 QEMU model, rails 120/250/112 ns, 51 GB/s · HW: Threadripper 7960X + Montage CXL node (node1) vs local DDR5 (node0), 64 B events
Model semantics: media write + one doorbell/GO per batch; HW roundtrip: NT store + sfence + read-back · median of repeated runs

(A) Notification latency vs. batch size



(B) Single-stream throughput



(C) Rail calibration check

